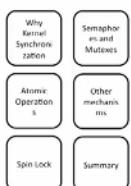


Kernel Synchronization



适用范围 / 情境.

Why sync? Overview

- Shared data
- Critical regions (critical sections)
 - Code paths that access/manipulate shared data

```

993 static int copy_fs(unsigned long clone_flags, struct task_struct *tsk)
994 {
995     struct fs_struct *fs = current->fs;
996     if (clone_flags & CLONE_FS)
997         /* tsk->fs is shared by what we want */
998         spin_lock(fs->lock);
999     if (fs->in_exec)
1000         spin_unlock(fs->lock);
1001     else
1002         spin_unlock(fs->lock);
1003     if (fs->users++ == 1)
1004         spin_lock(fs->lock);
1005     else
1006         spin_unlock(fs->lock);
1007     if (tsk->fs = copy_fs_struct(fs));
1008     if (!tsk->fs)
1009         return -ENOMEM;
1010     return 0;

```

(shared data)

Define "共享"? → critical region 里不包含或者 code
只可取少地包含操作 critical region 的程序.

fs → 本身不是 shared data

current → fs 是 shared data

critical region

操作 critical
data (操作)
锁且
foo-
程序及
data
操作

Overview (Cont.)

- Race condition
 - Two threads execute simultaneously within the same critical region
- Synchronization
 - Prevent unsafe concurrency and avoid race conditions

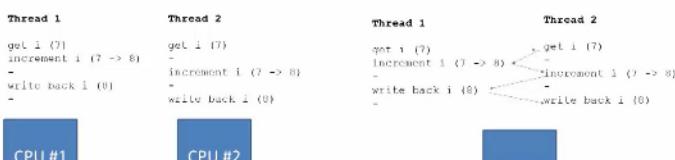
两个 thread 都在 critical section
里 (操作 shared data)
→ race condition
(不易重现 bug)

e.g. 两个 CPU 都在跑 Kernel
同一段操作 shared data
的程序

Example

- i: shared data
- i++

Thread 1	Thread 2
get i (7)	=
increment i (7 -> 8)	=
write back i (8)	=
=	get i (8)
=	increment i (8 -> 9)
=	write back i (9)

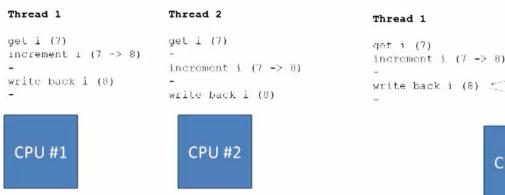


Example

競爭
競爭

- i: shared data
- i++

Thread 1	Thread 2
get i (7)	-
increment i (7 -> 8)	-
write back i (8)	-
-	get i (8)
-	increment i (8 -> 9)
-	write back i (9)



Source: Robert Love, "Linux Kernel Development"

用CPU的设计者提供

Solution

- Atomic instruction
 - X86 instruction, compare and exchange (CMPSXCHG)

Thread 1	Thread 2
increment i (7 -> 8)	-
-	increment i (8 -> 9)
or	
-	increment i (7 -> 8)
increment i (8 -> 9)	-

缺点：

单一的变量用 atomic operation 解决就可以了。

但是大部分的 shared variable 都是一个数据结构
而不是一个变量。

单核：OK!

多核 CPU 必须保证

两个 atomic instruction 不会同时被
执行。

Locking

Locking

- Using atomic instruction to handle shared variable is OK, how about data structure?

Thread 1	Thread 2
try to lock the queue	try to lock the queue
succeed: acquired lock	failed: waiting ...
access queue ...	waiting ...
unlock the queue	waiting ...
...	succeed: acquired lock
	access queue ...
	unlock the queue
	...

Lock: 操作一段共享的数据的时候，哪怕我就是操作其中的一个变量我也要将他整个数据结构给加锁
一般的数据结构里面都包含了一个锁变量。

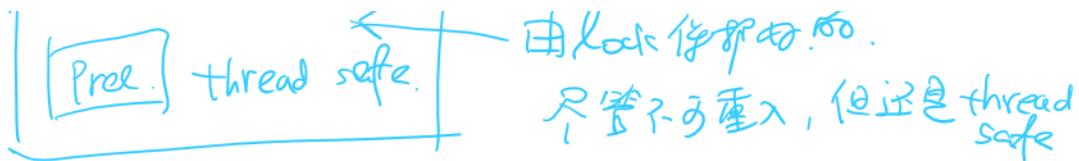
```
clone_flags & CLONE_FFS) {
    /* TSK->fs is already what we want */
    spin_lock(&fs->lock);
    if (fs->in_exec) {
        spin_unlock(&fs->lock);
        return -EAGAIN;
    }
    fs->users++;
    spin_lock(&fs->lock);
    return 0;
}
fs->users--;
spin_unlock(&fs->lock);
```

只是操作了 user，但也要
lock 整个数据结构。

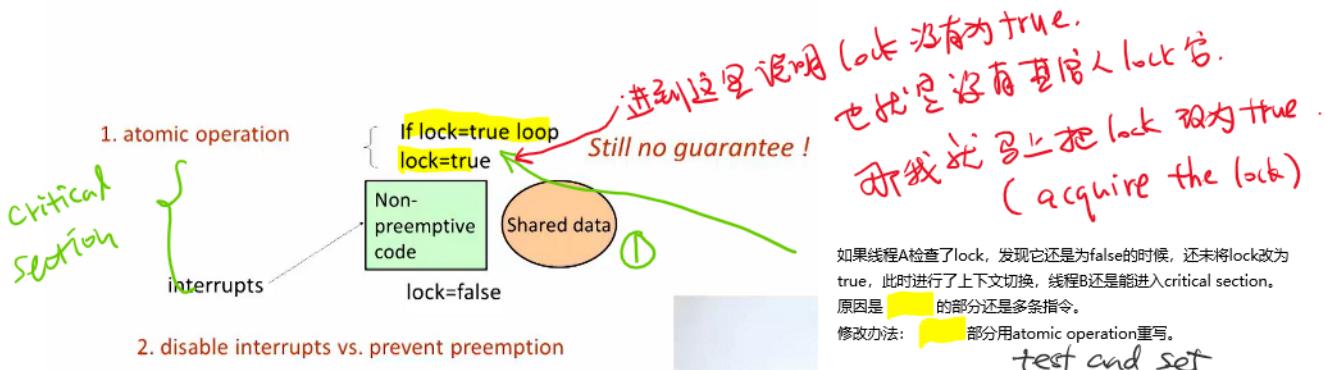
Preemption Vs. thread safe.

Pre. thread safe.

由 lock 保护的。
对谁，对谁，谁锁定谁



Locking



有2个问题：① ②

②. 就算进到 critical section 了也不能保证
不会出现 racing condition, 因为这个 thread
可能会被 interrupt 打断. 而且不能保证
interrupt subroutine 不会操作 shared data.

(如果能保证 interrupt subroutine 不会操作 shared data, 那么只要
在进入用 atomic operation 就能了)

interrupt subroutine 不操作 shared data, 但引发了 context switch
e.g. timer 会定期打进来, 进而执行 timer foo subroutine. 执行完后会
call scheduler, 而无法保证其他 foo thread 会不会 call 进 kernel
对 shared data 进行操作.

(解决办法是把 NEED_RESCHEDULE 关掉, Kernel 在进到
interrupt subroutine 时候会检查这个 flag, 如果关了就不会
call Scheduler)

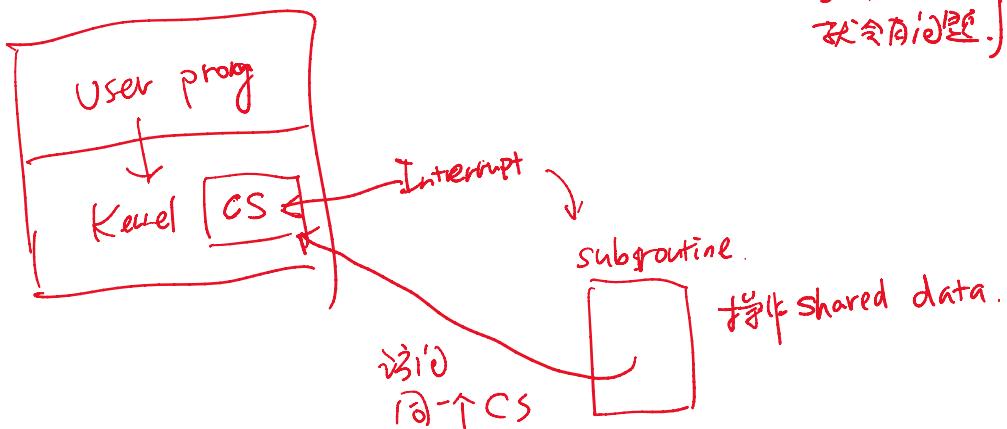
OS 支持 preemption:

OS 支持在某一时刻
突然停下来执行另一个
(thread) foo 程序,

- Pseudo concurrency (单核)
 - Two threads do not actually execute simultaneously but interleave with each other (in critical region)
- True concurrency (多核)
 - Two threads execute simultaneously on two processors (SMP)

Causes of Concurrency (Cont.)

- Interrupts (*Interrupt subroutine 抢占 shared data*)
 - An interrupt can occur asynchronously at almost any time, interrupting the currently executing code
- Softirqs and tasklets
 - The kernel can raise or schedule a softirq or tasklet at almost any time, interrupting the currently executing code
- Kernel preemption
 - Because the kernel is preemptive, one task in the kernel can preempt another (*用户态程序向内核发送 preemption, 但用户态程序同一块 shared data*)
- Sleeping and synchronization with user-space
 - A task in the kernel can sleep and thus invoke the scheduler, resulting in the running of a new process (*hold resource 在时候是不能 sleep 的。因为会进行 reschedule*)
- Symmetrical multiprocessing
 - Two or more processors can execute kernel code at exactly the same time



Causes of Concurrency (Cont.)

- Interrupt-safe
 - Code that is safe from concurrent access from an interrupt handler
- SMP-safe
 - Code that is safe from concurrency on symmetrical multiprocessing machines
- preempt-safe
 - Code that is safe from concurrency with kernel preemption

Kernel 支持 Interrupt \neq Preempt

Preempt:

一个进程在执行内核程序的时候，除非自愿放弃执行，否则另一个进程的用户态程序不能抢占执行。也就是说，内核是不可抢占或者说内核是不可重入的。

在 kernel 中工作时
一半，
e.g. 在 prog A 还没有 lock 任何 shared resource 时候就收到 interrupt.

Preemption Kernel 是什么？
定时器很高

Linux：早期仅支持 Interrupt. 后期才支持 preemption to SMP.

Protect the data structure !

Not the code !

因此不要 Lock 锁定的范围设定得过大.

BKL → 切切切 → 如今 too Kernel 不再被 interrupt too 区域已长得子了.

准时 Real Time < ①. 做事要快

②. CPU 什么时候把事情交给子

Improve Respond Time: (CPU 什么时候能“打进来”，Lock 粒度越细.

把 too critical section

Interrupt 越能打进来. Response Time 越快.

拆成一些极小的 critical section

实时性越高.)

Interrupt 打进来到退出的时间是在 1ms 以下，可以称之为 Real Time OS 3.

经常使用 Spin Lock 会让 Response Time 变大.

经常被打断：事情不能专心做 (实时性和耗时的 trade off)

Questions to ask before locking

绝大部分的 Kernel data structure 都不是 global 的.

- Is the data global? Can a thread of execution other than the current one access it? 是不是 per CPU 的?
- Is the data shared between process context and interrupt context? Is it shared between two different interrupt handlers? between processes
between process and int.
between int hand
- If a process is preempted while accessing this data, can the newly scheduled process access the same data?
- Can the current process sleep (block) on anything? If it does, in what state does that leave any shared data?

如果 Lock 了一个资源之后去睡觉，那可能会出问题，因为 Lock 的资源别人可能要用！

△ 操作系统设计出这么多种 Synchronization 机制的原因在于保护 share data 有不同的等级 (不同的问题)

Deadlocks

[prevent or resolve]

一般OS采取的策略.



- Racing Condition
- Dead Lock .

Sleep waiting or Busy waiting forever !

How to prevent?

- Atomicity vs. ordering ← request resource follows order and atomic
- Make sure lock orders
 - Implement lock ordering
 - Nested locks must always be obtained in the same order
 - Do not double acquire the same lock
 - The order of unlock does not matter with respect to deadlock

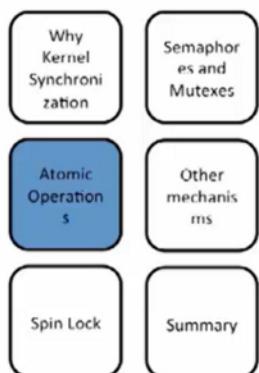
particular

Lock Req: A → B → C

→ 一般来说

Unlock Req: C → B → A

解锁要求: CBA



Linux 同步机制之一 :

Atomic Operations

— 提供了一些不会被打断的指令

Atomic Operations

- Atomic integer operations (+ -)
- Atomic bitwise operations ($0 \Leftrightarrow 1$)

Atomic operation 是 Hardware dependent 的。
不同 Hardware Platform 都有自己的实现方式
(不同 Atomic 指令)

How to protect critical section?

- Atomic Integer Operations

ATOMIC_INIT(int i)	At declaration, initialize an atomic_t to i
int atomic_read(atomic_t *v)	Atomically read the integer value of v
void atomic_set(atomic_t *v, int i)	Atomically set v equal to i
void atomic_add(int i, atomic_t *v)	Atomically add i to v
void atomic_sub(int i, atomic_t *v)	Atomically subtract i from v
void atomic_inc(atomic_t *v)	Atomically add one to v
void atomic_dec(atomic_t *v)	Atomically subtract one from v
int atomic_sub_and_test(int i, atomic_t *v)	Atomically subtract i from v and return true if the result is zero; otherwise false
int atomic_add_negative(int i, atomic_t *v)	Atomically add i to v and return true if the result is negative; otherwise false
int atomic_dec_and_test(atomic_t *v)	Atomically decrement v by one and return true if zero; false otherwise
int atomic_inc_and_test(atomic_t *v)	Atomically increment v by one and return true if the result is zero; false otherwise

保护一个 Integer 变量。

Atomic operation

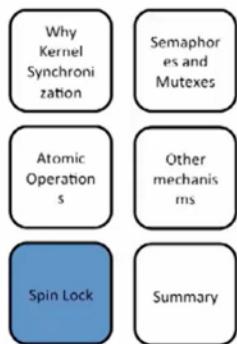
- Atomic Bitwise Operations

void set_bit(int nr, void *addr)	Atomically set the nr-th bit starting from addr
void clear_bit(int nr, void *addr)	Atomically clear the nr-th bit starting from addr
void change_bit(int nr, void *addr)	Atomically flip the value of the nr-th bit starting from addr
int test_and_set_bit(int nr, void *addr)	Atomically set the nr-th bit starting from addr and return the previous value
int test_and_clear_bit(int nr, void *addr)	Atomically clear the nr-th bit starting from addr and return the previous value
int test_and_change_bit(int nr, void *addr)	Atomically flip the nr-th bit starting from addr and return the previous value
int test_bit(int nr, void *addr)	Atomically return the value of the nr-th bit starting from addr

get

Linux里所有的保护机制都是基于
Atomic operation 来实现的。

e.g. Spin Lock 的实现。



Linux 同步机制之二：

Spin Lock

[适用范围]
— 短时间的 Busy Waiting. SMP.

Spin lock

- spin_lock

- Normally, critical section is more than just a variable, spin lock protects a sequence of codes manipulating the critical section
- Perform busy loops for waiting the lock to become available
- Spin_lock disables the kernel preemption
- On UP the locks compile away and do not exist; they simply act as markers to disable and enable kernel preemption
- If kernel preempt is turned off, the locks compile away entirely

不希望有人抢占打断 (因为 get the resource)

不希望又有人进来操纵这个变量
不会 reentry.

单核
Uni-processor

spin_lock implementation on SMP

关闭 preemption

```
1 #define BUILD_LOCK_OPS(op, locktype)
2 void __lockfunc __raw_##op##_lock(locktype##_t *lock)
3 {
4     for (;;) {
5         preempt_disable();
6         if (!likely(__raw_##op##_trylock(lock)))
7             break;
8         preempt_enable();
9
10        if (!(lock->lock))
11            (lock->lock) = 1;
12        while (!__raw_##op##_can_lock(lock) && (lock->lock))
13            arch_##op##_relax(&lock->raw_lock);
14    }
15    (lock->lock) = 0;
16 }
```

try to lock (atomic operation)

没获得 Lock, 再把 preemption 打开, 再次 spin.

(说明另外一个CPU在critical section中)

为向量是在等待另一个 thread 退出 critical section 的时候, 把此进线程放到低功耗模式下.

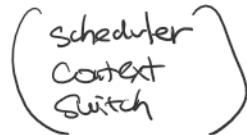
△ 注意, 如果另一个CPU永远不退出 CS, 那么执行 spinlock 从这颗 CPU 就“死掉”了.

- 时候，把此进线放到低功耗模式下。
- △ 注意，如果另一个CPU永远不退出CS，那么执行spin lock的这颗CPU就“死掉”了。
 - △ 如果要lock很久，不建议使用spin lock。

Spin lock

- Principles for using spin lock
 - Lock the data, not the code
 - Don't hold a spin lock for a long time (because someone may wait outside)
 - Don't recursively spin lock in Linux (like you lock your key in your car)
 - Spin lock (instead of using mutex) if you can finish within 2 context switches time
 - In a preemptive kernel, spin lock will perform preempt_disable() 非local CPU不会 interrupt
 - Can be used in interrupt handler (mutex/semaphores cannot) must disable "local" interrupt
- 别人等很久，那么这些时候与其做spin，不如做些其他事情。
(其他人忙于Busy waiting)
- spin lock是圆 spin lock。
一定得等到从局部 spin，从而外面的
spin lock永远都不解 unlock。
- Spin lock是希望不进行 Context Switch 等到别人 release lock 才自己进入CS。
而如果你的事情太长，完全可以先 context switch 去做其他的事，等做完了再 context switch 回来。(这样还不会浪费CPU)
如果等的时间短，Context switch是不经济的。 (Spin 申请 < Context Switch)
开销

★ 不被抢占，但可以被 interrupt。



如果 data is shared by Process and interrupt handler.
则需要手动关闭 interrupt handler.

Spin lock

- When you use spin_lock(x) in the kernel, make sure you don't spend too much time in the lock, and there is no recursive spin_lock(x)
 - No spin_lock(x) in the spin_lock(x) period
- When you use spin_lock(x) in interrupt handler, make sure the interrupt has been disabled. Otherwise, use spin_lock_irq(x) or spin_lock_irqsave(x)

= spinlock + disable interrupt

加强版 spin_lock_irq()

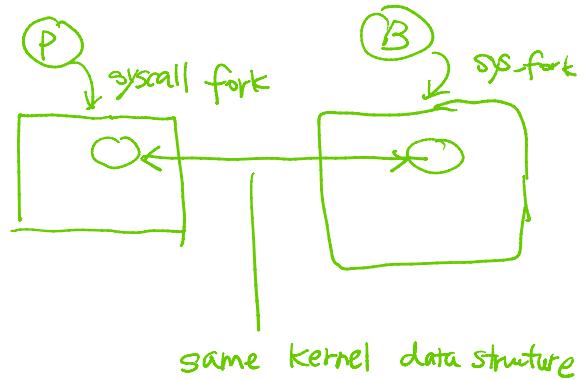
Considering following situations

1. You are writing kernel code (process context)
 - Shared data between processes
 2. You are writing interrupt handler (interrupt context)
 - Shared data between interrupt handler
 3. You are writing kernel code or interrupt handler
 - Shared data between interrupt handlers and processes
- How could you use spin_lock to protect shared data?

1. 3以

2. 基本的 spin_lock 不行. 需要先 disable interrupt.

3. 同 2.



Non-preemption

OK!

Interrupt

Not safe

Spin lock irq

- Disables interrupts and acquires the lock
- spin_unlock_irq() unlocks the given lock
- Disable interrupts (spin_lock_irq) and enable interrupts anyway (spin_lock_irqsave) no matter the status of current interrupts

– Use of spin_lock_irq() therefore is not

recommended 很少用. 因为无论发生 interrupt 是开启. 都会进行 disable&enable.
如果发生 interrupt 是 disable 了. 那么行为就会有差异.

Spin lock irqsave

- Disables interrupt if it is not disabled and acquires the lock
- spin_unlock_irqrestore() unlocks the given lock and returns interrupts to their previous state

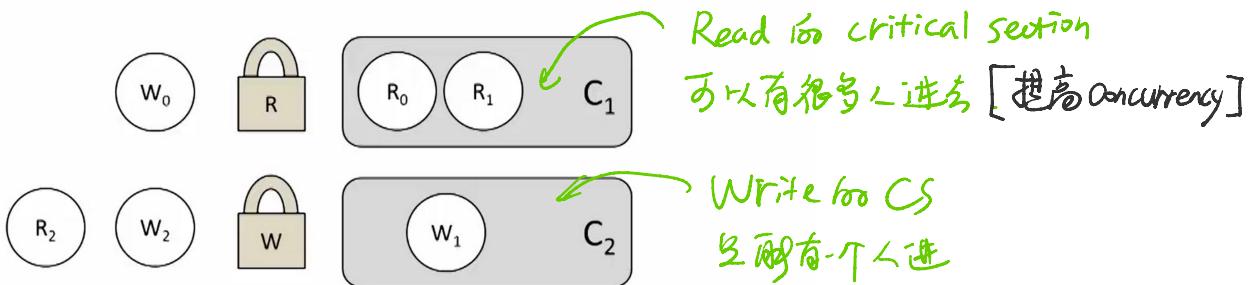
慢. (save and restore interrupt 状态)
但是保险.

Spin lock methods

spin_lock()	Acquires given lock
spin_lock_irq()	Disables local interrupts and acquires given lock
spin_lock_irqsave()	Saves current state of local interrupts, disables local interrupts, and acquires given lock
spin_unlock()	Releases given lock
spin_unlock_irq()	Releases given lock and enables local interrupts
spin_unlock_irqrestore()	Releases given lock and restores local interrupts to given previous state
spin_lock_init()	Dynamically initializes given spinlock_t
spin_trylock()	Tries to acquire given lock; if unavailable, returns nonzero
spin_is_locked()	Returns nonzero if the given lock is currently acquired, otherwise it returns zero

Kernel 中 foo DS 有些是一个写，有很多人可以读后，这种用 spinlock 不是很效率。

Read/Write Spin Locks



Read write spin lock

- Lock usage can be clearly divided into readers and writers (e.g. search and update)
- Linux task list is protected by a reader-writer spin lock

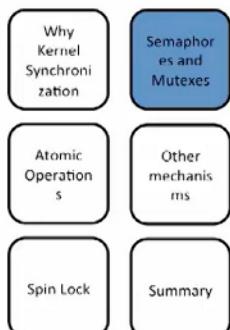
```
read_lock(&mr_rwlock);
/* critical section (read only) ... */ /* critical section (read and write) ... */
read_unlock(&mr_rwlock);
write_lock(&mr_rwlock);
write_unlock(&mr_rwlock);
```

- Must split into reader/writer paths
 - You cannot do this
- ```
read_lock(&mr_rwlock);
write_lock(&mr_rwlock);
```



# Read write spin lock

| Method                    | Description                                                                                                     |
|---------------------------|-----------------------------------------------------------------------------------------------------------------|
| read_lock()               | Acquires given lock for reading                                                                                 |
| read_lock_irq()           | Disables local interrupts and acquires given lock for reading                                                   |
| read_lock_irqsave()       | Saves the current state of local interrupts, disables local interrupts, and acquires the given lock for reading |
| read_unlock()             | Releases given lock for reading                                                                                 |
| read_unlock_irq()         | Releases given lock and enables local interrupts                                                                |
| read_unlock_irqrestore()  | Releases given lock and restores local interrupts to the given previous state                                   |
| write_lock()              | Acquires given lock for writing                                                                                 |
| write_lock_irq()          | Disables local interrupts and acquires the given lock for writing                                               |
| write_lock_irqsave()      | Saves current state of local interrupts, disables local interrupts, and acquires the given lock for writing     |
| write_unlock()            | Releases given lock                                                                                             |
| write_unlock_irq()        | Releases given lock and enables local interrupts                                                                |
| write_unlock_irqrestore() | Releases given lock and restores local interrupts to given previous state                                       |
| write_trylock()           | Tries to acquire given lock for writing; if unavailable, returns nonzero                                        |
| rw_lock_init()            | Initializes given rwlock_t                                                                                      |
| rw_is_locked()            | Returns nonzero if the given lock is currently acquired, or else it returns zero                                |



Linux 同步机制之三：

Semaphores and Mutexes

——适用于需要很久的

- Mutex is a special case semaphore!

## Semaphores

- sleeping locks
- do not disable kernel preemption
- the mutex/semaphore place the task onto a wait queue and put the task to sleep
- well suited to locks that are held for a long time
- not optimal for locks that are held for very short periods
- can be obtained only in process context, **not interrupt context**
- can sleep while holding a semaphore
- **cannot hold a spin lock while you acquire a semaphore**
  - How about hold a semaphore while you acquire a spin lock

必须快速做完  
(因为在 interrupt 才干进事)  
By

因为 semaphore 会把 process put into sleep. 而且不能同时 hold 3 -> 1 spin lock.  
所以持有的时候不能 hold 3 spin lock 并做了 sleep - 这是绝对不行的！

反过来，已~~hold~~ hold Semaphore. 可以 hold spinlock.

## Semaphores (Cont.)

- binary semaphore vs. counting semaphore  
*(not used much in the kernel)*
- down\_interruptible()
  - attempts to acquire the given semaphore. If it fails, it sleeps in the TASK\_INTERRUPTIBLE
- down()
  - If it fails, places the task in the TASK\_UNINTERRUPTIBLE state if it sleeps  
*(normally not used)*

## Semaphores

|                                        |                                                                                                |
|----------------------------------------|------------------------------------------------------------------------------------------------|
| sema_init(struct semaphore *, int)     | Initializes the dynamically created semaphore to the given count                               |
| init_MUTEX(struct semaphore *)         | Initializes the dynamically created semaphore with a count of one                              |
| init_MUTEX_LOCKED(struct semaphore *)  | Initializes the dynamically created semaphore with a count of zero (so it is initially locked) |
| down_interruptible(struct semaphore *) | Tries to acquire the given semaphore and enter interruptible sleep if it is contended          |
| down(struct semaphore *)               | Tries to acquire the given semaphore and enter uninterruptible sleep if it is contended        |
| down_trylock(struct semaphore *)       | Tries to acquire the given semaphore and immediately return nonzero if it is contended         |
| up(struct semaphore *)                 | Releases the given semaphore and wakes a waiting task, if any                                  |

down(—)

立马返回

## Implementation

```
1 void down(struct semaphore *sem)
2 {
3 unsigned long flags;
4 raw_spin_lock_irqsave(&sem->lock, flags);
5 if (likely(sem->count > 0))
6 sem->count--;
7 else
8 __down(sem);
9 raw_spin_unlock_irqrestore(&sem->lock, flags);
10 }
11 }
```

拿到锁后，Semaphore->lock之后再对 Semaphore (spinlock)

```

1 static noinline void __sched __down(struct semaphore *sem)
2 {
3 __down_common(sem, TASK_UNINTERRUPTIBLE, MAX_SCHEDULE_TIMEOUT);
4 }
5 static inline int __sched __down_common(struct semaphore *sem, long s:
6 long tim
7 {
8 struct task_struct *task = current;
9 struct semaphore_waiter waiter;
10 list_add_tail(&waiter.list, &sem->wait_list);
11 waiter.task = task;
12 waiter.up = 0;
13 for (;;) {
14 if (signal_pending_state(state, task))
15 goto interrupted;
16 if (timeout <= 0)
17 goto timed_out;
18 __set_task_state(task, state);
19 raw_spin_unlock_irq(&sem->lock);
20 timeout = schedule_timeout(timeout);
21 raw_spin_lock_irq(&sem->lock);
22 if (waiter.up)
23 return 0;
24 }
25 timed_out:
26 list_del(&waiter.list);
27 return -ETIME;
28 interrupted:
29 list_del(&waiter.list);
30 return -EINTR;
31 }
32 }

```

TASK\_UNINTERRUPTIBLE

是 Idle mode , Scheduler 不会 schedule 31.

Linux 不会 schedule 31 Busy too .



## Mutex

- Behaves similar to a semaphore with a count of one, but it has a simpler interface, more efficient performance

```

mutex_lock(&mutex);
/* critical region ... */
mutex_unlock(&mutex);

```

| Method                           | Description                                                                                             |
|----------------------------------|---------------------------------------------------------------------------------------------------------|
| mutex_lock(struct mutex *)       | Locks the given mutex; sleeps if the lock is unavailable                                                |
| mutex_unlock(struct mutex *)     | Unlocks the given mutex                                                                                 |
| mutex_trylock(struct mutex *)    | Tries to acquire the given mutex; returns one if successful and the lock is acquired and zero otherwise |
| mutex_is_locked (struct mutex *) | Returns one if the lock is locked and zero otherwise                                                    |

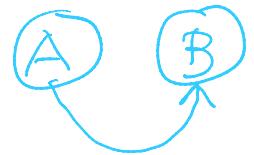
只锁住  
解锁 .



# Other Mechanisms

- Completion variables

- an easy way to synchronize between two tasks in the kernel when one task needs to signal to the other that an event has occurred



- Sequential locks

- shortened to seq lock, is a newer type of lock introduced in the 2.6 kernel. It provides a simple mechanism for reading and writing shared data

## Other Mechanisms (Cont.)

- Interrupt disable

- Allows a kernel control path to continue executing even when hardware issues IRQ signals

- Preemptions disable

- For per-processor data access. If no spin locks are held, the kernel is preemptive, and it would be possible for a newly scheduled task to access this same variable

## Other Mechanisms (Cont.)

- Memory barriers

- May order

a = 1;  
b = 2;

- Will not reorder

a = 1;  
b = a; *dependency*

- Example, a=1, b=2, no mb(), rmb()

*Read Memory Barrier*

Thread 1

a = 3;

mb();

b = 4;

—

—

Thread 2

—

—

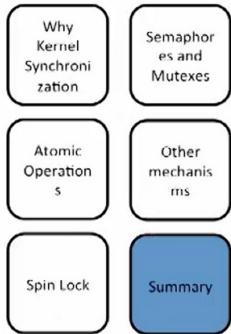
c = b;

Kern...mb();

Synchro...

d = a;

开始录音: 2019年11月24日 19:10



## Protection of Shared Data in Kernel

| Requirement                         | Recommended Lock              |
|-------------------------------------|-------------------------------|
| Low overhead locking                | Spin lock                     |
| Short lock hold time                | Spin lock                     |
| Long lock hold time                 | Mutex / Sem                   |
| Need to lock from interrupt context | Spin lock & disable interrupt |
| Need to sleep while holding lock    | Mutex                         |

## Protection of Shared Data in Kernel

| Kernel control paths accessing the data structure | UP protection             | MP further protection |
|---------------------------------------------------|---------------------------|-----------------------|
| Exceptions                                        | Semaphore                 | None                  |
| Interrupts                                        | Local interrupt disabling | Spin lock             |
| Deferrable functions                              | None                      | None or spin lock     |
| Exceptions + Interrupts                           | Local interrupt disabling | Spin lock             |
| Exceptions + Deferrable functions                 | Local softirq disabling   | Spin lock             |
| Interrupts + Deferrable functions                 | Local interrupt disabling | Spin lock             |
| Exceptions + Interrupts + Deferrable functions    | Local interrupt disabling | Spin lock             |

*Will discuss deferrable functions (softirq and tasklet latter)*

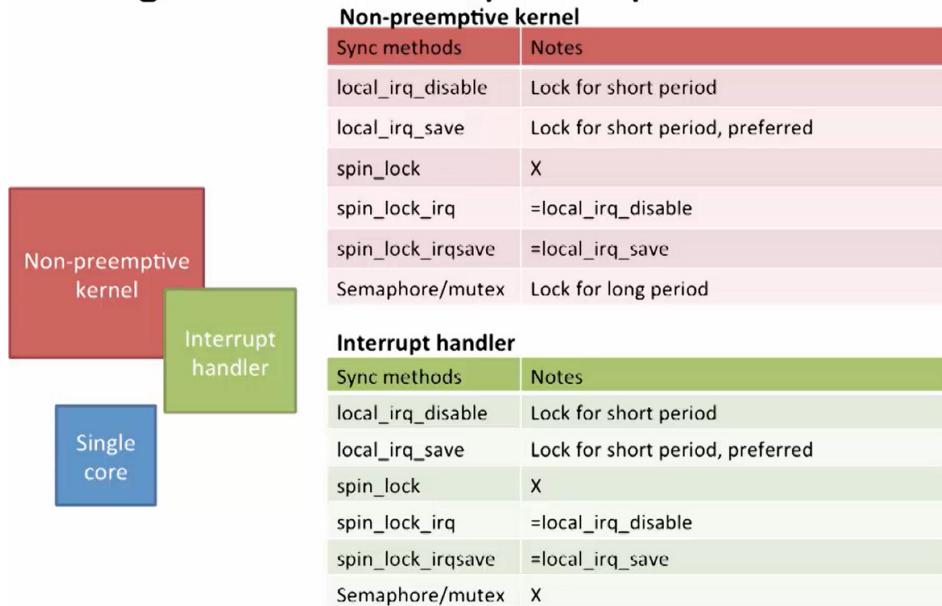
*tasklet & Soft IRQ.*

## Protection of Shared Data in Kernel

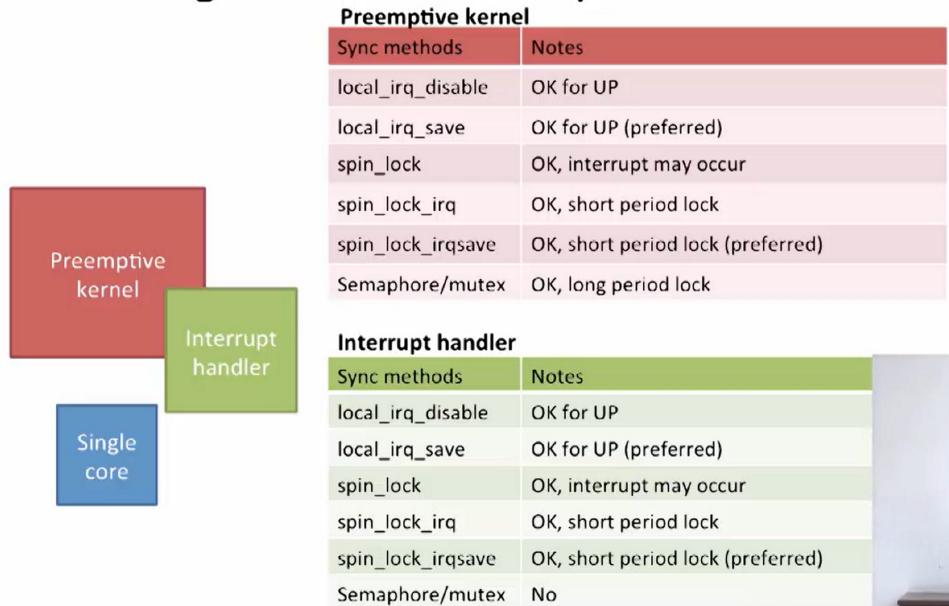
- Exception handler
  - Such as system call, data can be shared by processes

Single Core + Non-preemptive kernel

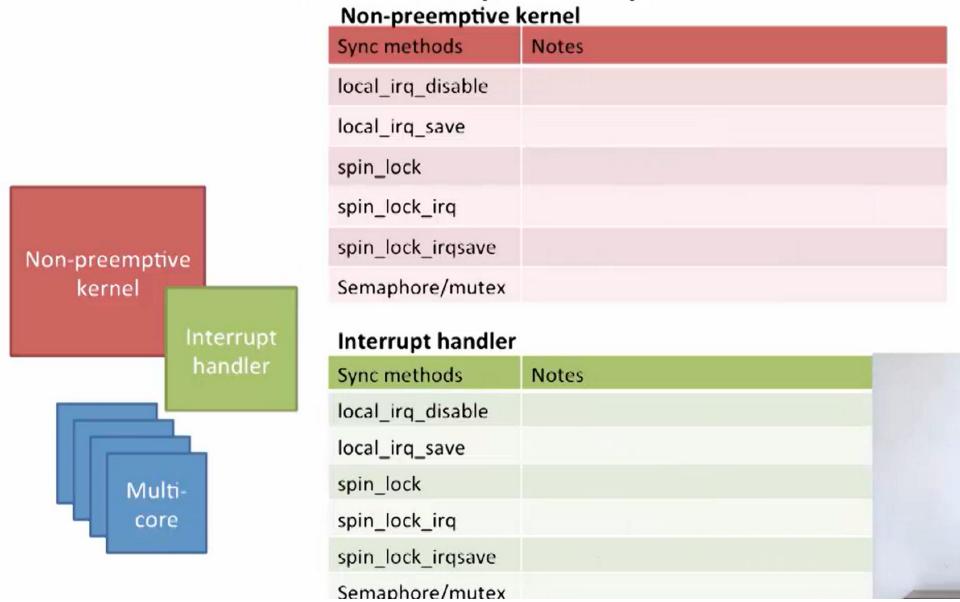
# Single Core + Non-preemptive kernel



# Single Core + Preemptive kernel



## Multi-Core + Non-preemptive kernel



## Multi-Core + Preemptive kernel

